

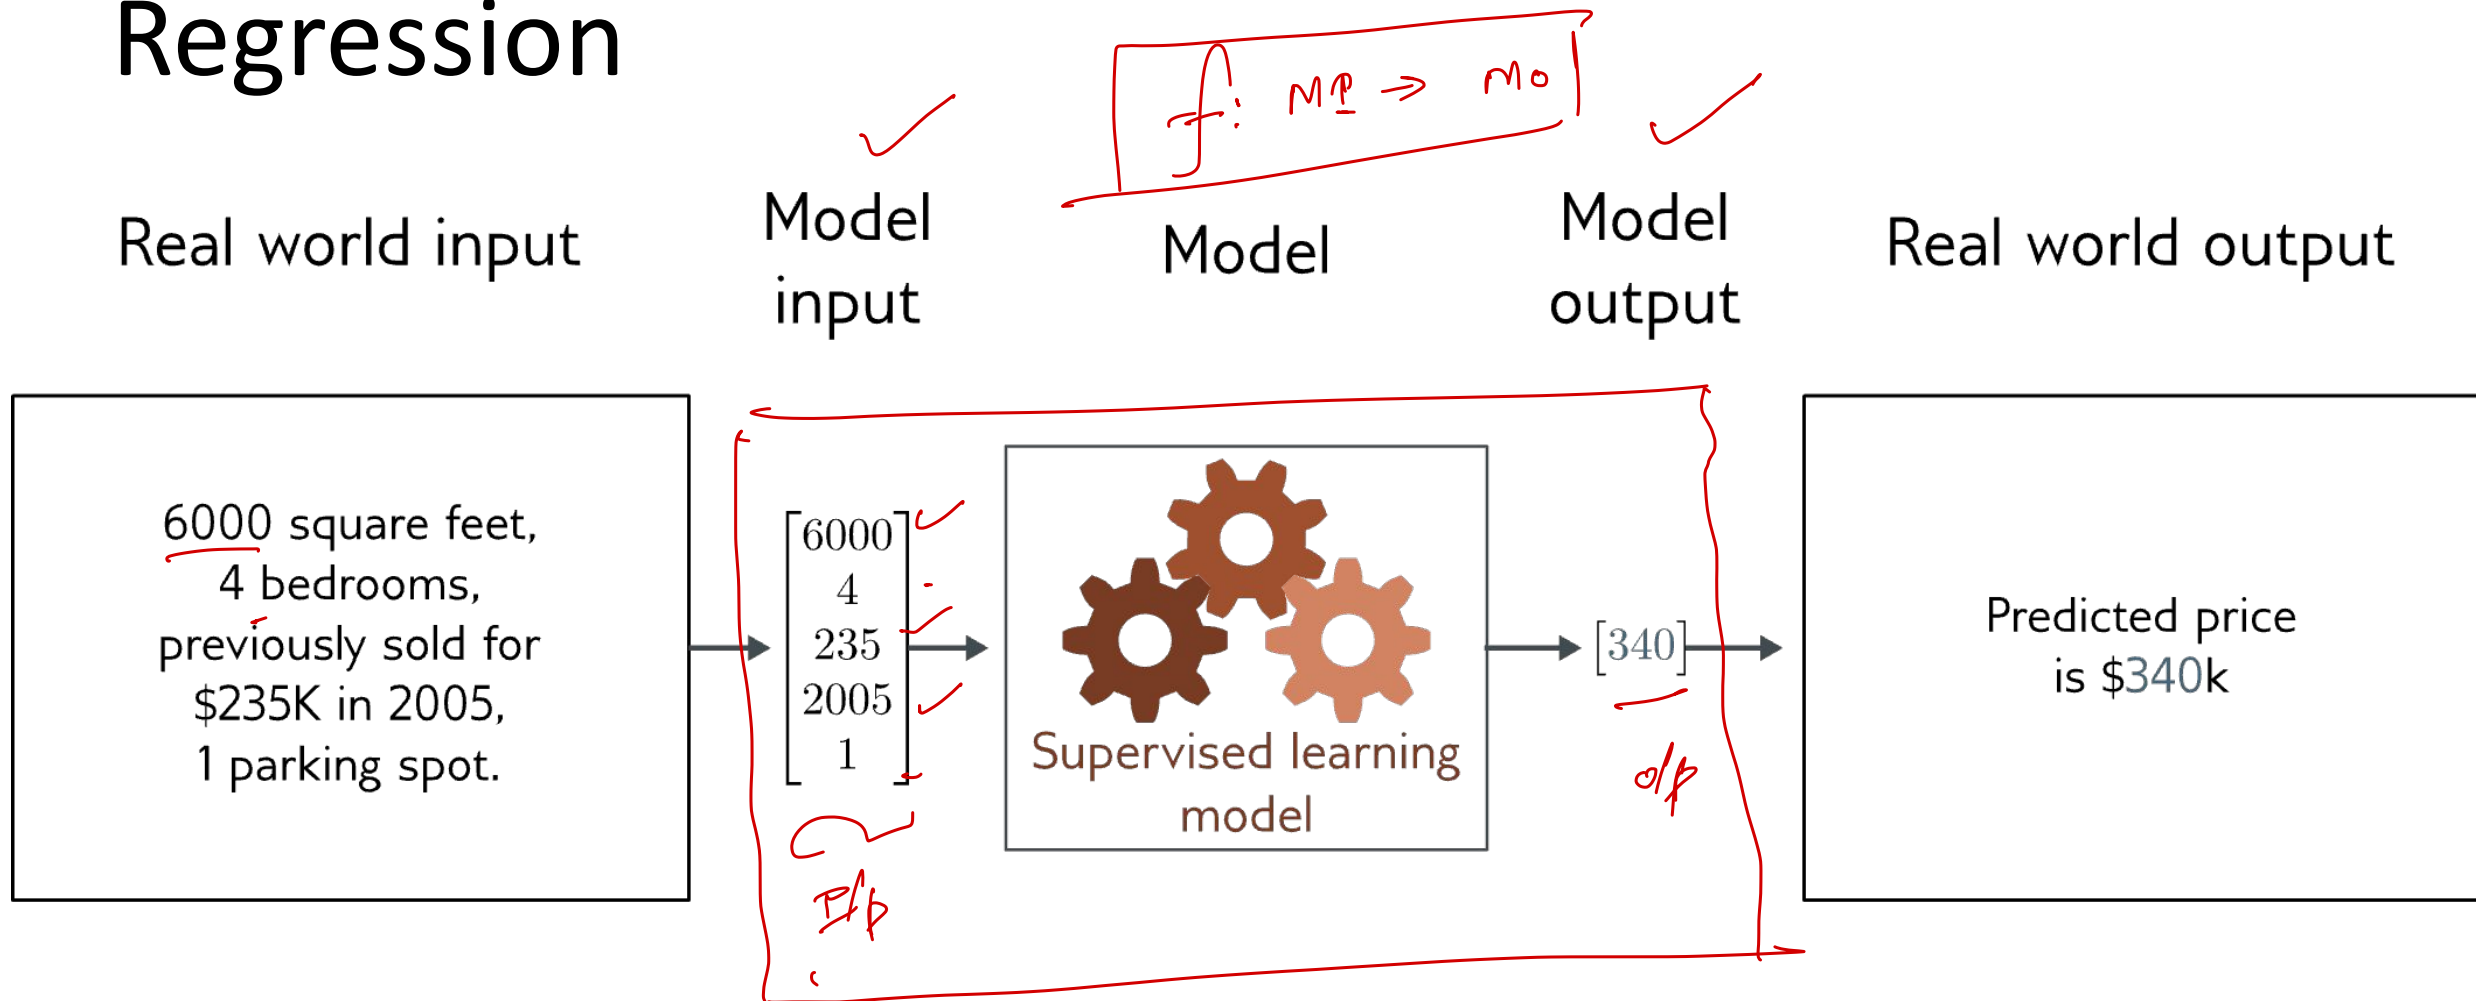
Supervised Learning

January 3rd, 2025

Deep Learning (CS60010)

**Slides adapted from <http://udlbook.com>*

Regression



- Univariate regression problem (one output, real value)

Supervised learning overview

- **Supervised learning model** = mapping from one or more inputs to one or more outputs
- Model is a mathematical equation
- Example:
 - Input is age and mileage of second hand Toyota Prius
 - Output is estimated price of car

Supervised learning overview

- **Supervised learning model** = mapping from one or more inputs to one or more outputs
- Model is a mathematical equation
- Computing the outputs from the inputs = **inference**
- Model also includes **parameters**
- Parameters affect outcome of equation

Training

↗

(Your model is completely known)

Parameters

in terms of parameters

$$y = w_0 + w_1 x$$

Supervised learning overview

Model
↓
Parameters
(use training)

- Supervised learning model = mapping from one or more inputs to one or more outputs

- ~~Model is a mathematical equation~~

- ✓ • Model is a family of equations

one partⁿ exⁿ [Training]

- Computing the outputs from the inputs = inference

$$y = w_0 + w_1 x$$

- Model also includes parameters

$$y = 1 + 2x$$
$$0.1 + 0.3x$$

- Parameters affect outcome of equation

- Training a model = finding parameters that predict outputs “well” from inputs for a training dataset of input/output pairs

$$0.7 - 1.4x$$

Notation:

- Input:

x

x
~~*x*~~
~~*x*~~

- Output:

y

- Model:

y = f[x]

Variables always Roman letters

Normal = scalar

Bold = vector

Capital Bold = matrix

Functions always square brackets

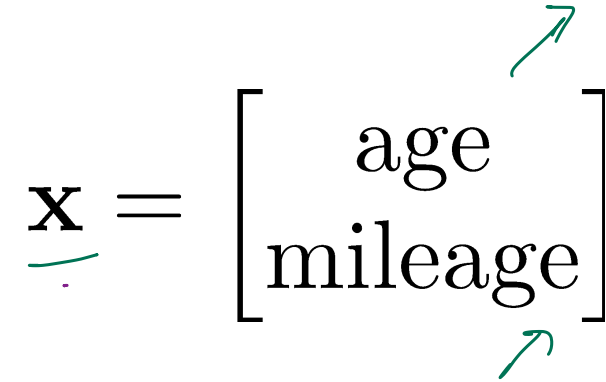
Normal = returns scalar

Bold = returns vector

Capital Bold = returns matrix

Notation example:

- Input:

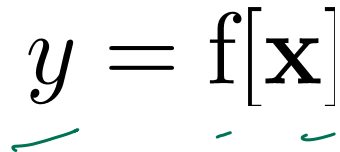
$$\underline{\mathbf{x}} = \begin{bmatrix} \text{age} \\ \text{mileage} \end{bmatrix}$$
The input vector \mathbf{x} is shown as a column vector containing 'age' and 'mileage'. The vector symbol \mathbf{x} is underlined in green. Green arrows point to the 'age' and 'mileage' elements. A red arrow points from the text 'Structured or tabular data' to the vector.

Structured or
tabular data

- Output:


$$y = [\text{price}] \rightarrow \text{scalar}$$
The output scalar y is shown as a single element in a vector notation [price]. The variable y is underlined in green. A green arrow points from the [price] to the handwritten word 'scalar'.

- Model:

$$\underline{y} = f[\underline{\mathbf{x}}]$$
The model equation $y = f[\mathbf{x}]$ is shown. The y is underlined in green, and the \mathbf{x} is underlined in green.

Model

- Parameters:



Parameters always Greek letters

- Model :

$$y = f[x, \phi]$$

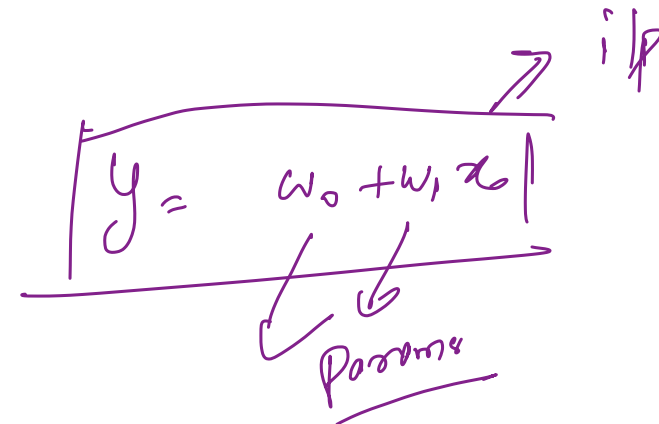
output

input

parameters

$$f = \frac{w_0 + w_1 x}{}$$

$$\phi = \begin{bmatrix} w_0 \\ w_1 \end{bmatrix}$$


$$y = w_0 + w_1 x$$

Loss function

- Training dataset of I pairs of input/output examples:

$$\{\mathbf{x}_i, \mathbf{y}_i\}_{i=1}^I$$

- Loss function or cost function measures how bad model is:

we define a loss fn \swarrow

$$L \left[\underbrace{\phi, f[\mathbf{x}, \phi]}_{\text{model}} , \underbrace{\{\mathbf{x}_i, \mathbf{y}_i\}_{i=1}^I}_{\text{train data}} \right]$$

I i/p o/p pairs

or for short:

$$L[\phi]$$

Returns a scalar that is smaller when model maps inputs to outputs better ✓

Training

- Loss function:

$$\underline{L[\phi]}$$

Returns a scalar that is smaller when model maps inputs to outputs better

- Find the parameters that minimize the loss:

$$\hat{\phi} = \operatorname{argmin}_{\phi} [L[\phi]]$$

Testing

- To test the model, run on a separate test dataset of input / output pairs

- See how well it generalizes to new data

X

↓

↑

don't show
the o/p

↗

Example: 1D Linear regression model

- Model:

$$\begin{aligned} \underline{y} &= f[x, \phi] \\ &= \underline{\phi_0} + \underline{\phi_1} \underline{x} \end{aligned}$$

- Parameters

$$\phi = \begin{bmatrix} \phi_0 \\ \phi_1 \end{bmatrix}$$

← y-offset ✓
← slope ✓

Example: 1D Linear regression model

- Model:


$$\begin{aligned} y &= f[x, \phi] \\ &= \phi_0 + \phi_1 x \end{aligned}$$

- Parameters

$$\phi = \begin{bmatrix} \phi_0 \\ \phi_1 \end{bmatrix}$$

← y-offset

← slope



Example: 1D Linear regression model

- Model:


$$\begin{aligned}y &= f[x, \phi] \\ &= \phi_0 + \phi_1 x\end{aligned}$$

- Parameters

$$\phi = \begin{bmatrix} \phi_0 \\ \phi_1 \end{bmatrix}$$

← y-offset

← slope



Example: 1D Linear regression model

- Model:


$$\begin{aligned} y &= f[x, \phi] \\ &= \phi_0 + \phi_1 x \end{aligned}$$

- Parameters


$$\phi = \begin{bmatrix} \phi_0 \\ \phi_1 \end{bmatrix}$$

← y-offset


← slope



Example: 1D Linear regression training data



Example: 1D Linear regression training data




Loss function:

$$\begin{aligned} \underline{L[\phi]} &= \sum_{i=1}^I (f[x_i, \phi] - y_i)^2 \\ &= \sum_{i=1}^I (\phi_0 + \phi_1 x_i - y_i)^2 \end{aligned}$$

“Least squares loss function”

Example: 1D Linear regression loss function




Loss function:

$$L[\phi] = \sum_{i=1}^I (f[x_i, \phi] - y_i)^2$$
$$= \sum_{i=1}^I (\phi_0 + \phi_1 x_i - y_i)^2$$

“Least squares loss function”

Example: 1D Linear regression loss function




Loss function:

$$L[\phi] = \sum_{i=1}^I (f[x_i, \phi] - y_i)^2$$
$$= \sum_{i=1}^I (\phi_0 + \phi_1 x_i - y_i)^2$$

“Least squares loss function”

Example: 1D Linear regression loss function




Loss function:

$$L[\phi] = \sum_{i=1}^I (f[x_i, \phi] - y_i)^2$$
$$= \sum_{i=1}^I (\phi_0 + \phi_1 x_i - y_i)^2$$

“Least squares loss function”

Example: 1D Linear regression loss function





Loss function:


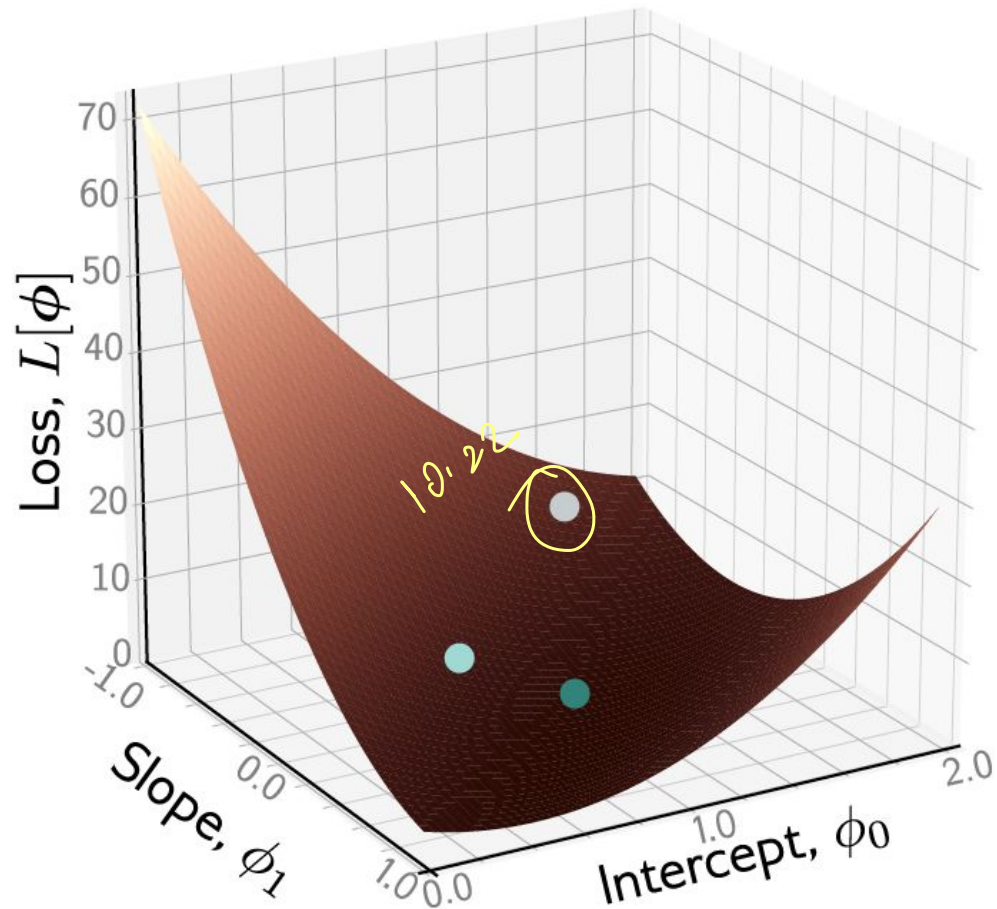
$$\begin{aligned} L[\phi] &= \sum_{i=1}^I (f[x_i, \phi] - y_i)^2 \\ &= \sum_{i=1}^I (\phi_0 + \phi_1 x_i - y_i)^2 \end{aligned}$$

“Least squares loss function”



Example: 1D Linear regression loss function



Example: 1D Linear regression loss function




Example: 1D Linear regression loss function




Example: 1D Linear regression loss function

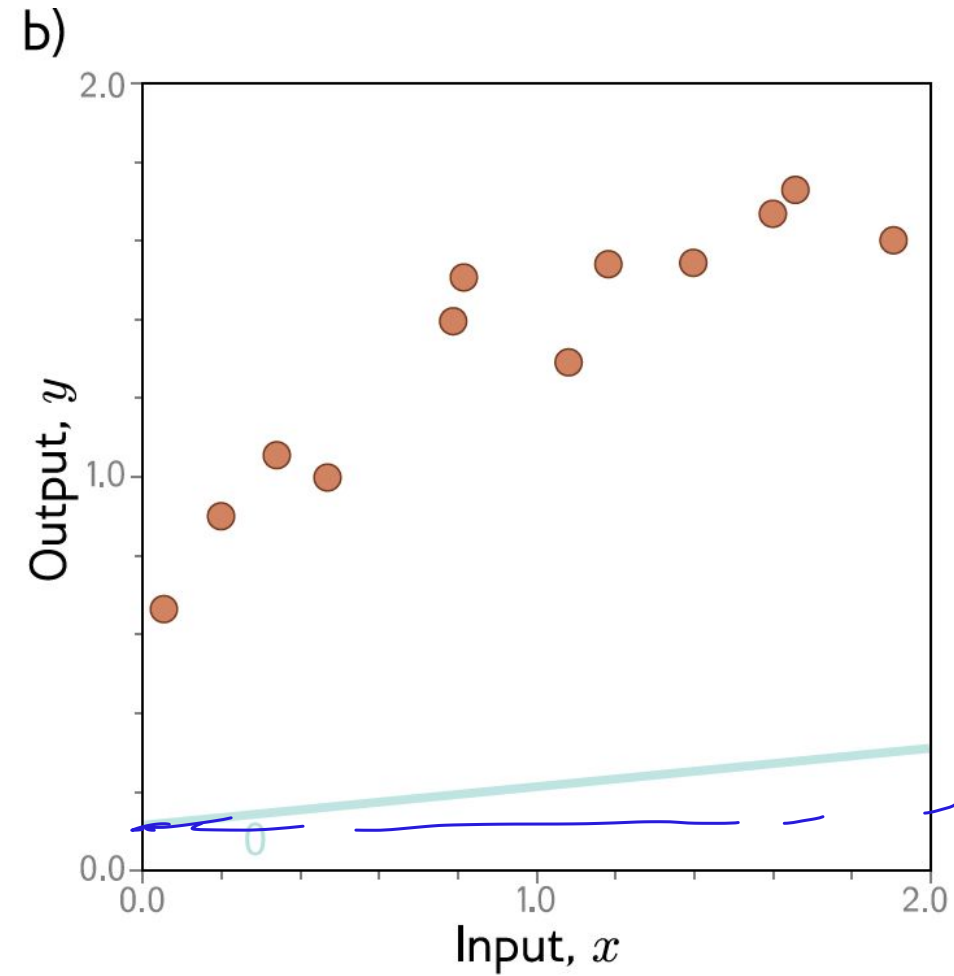

a)



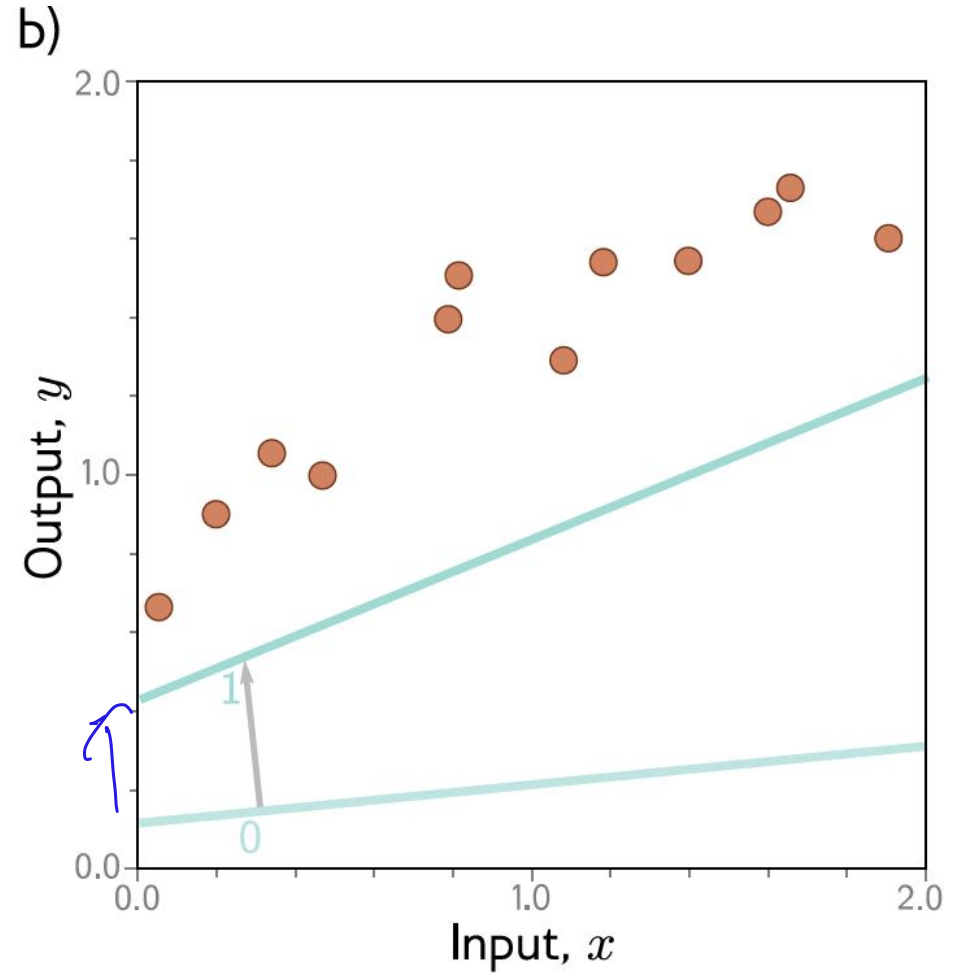

b)



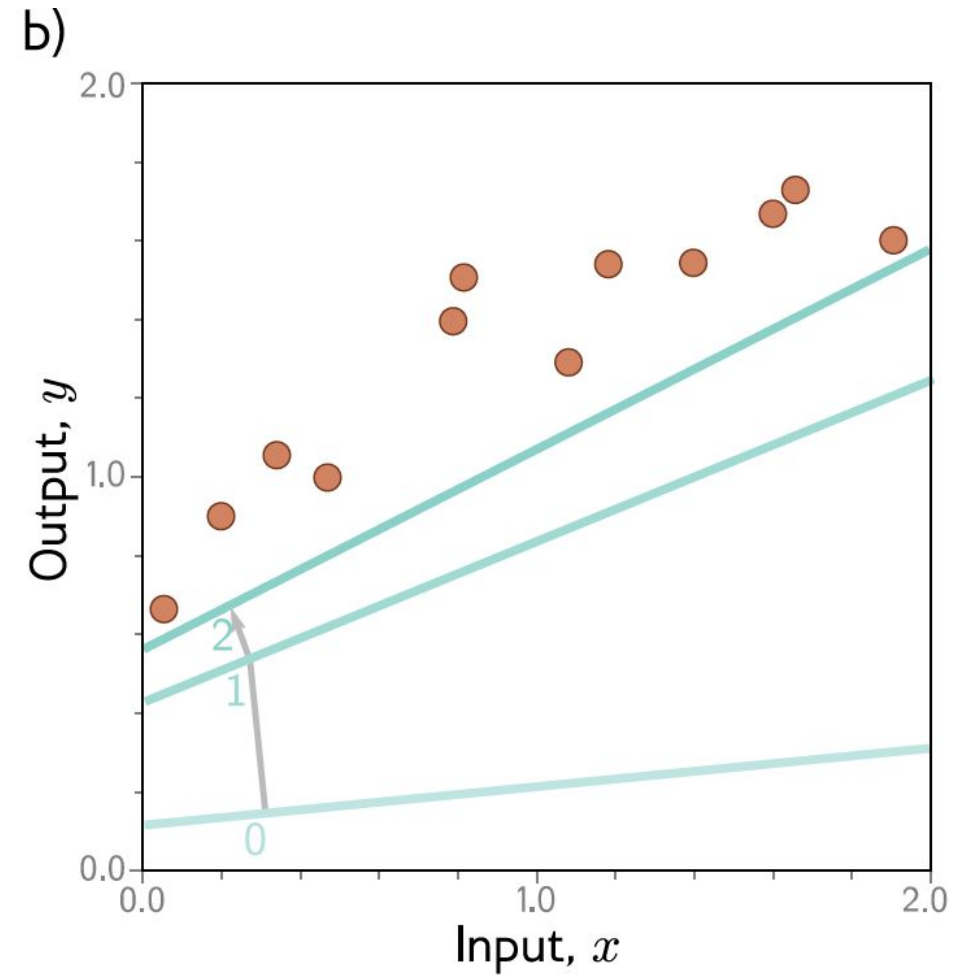

Example: 1D Linear regression training



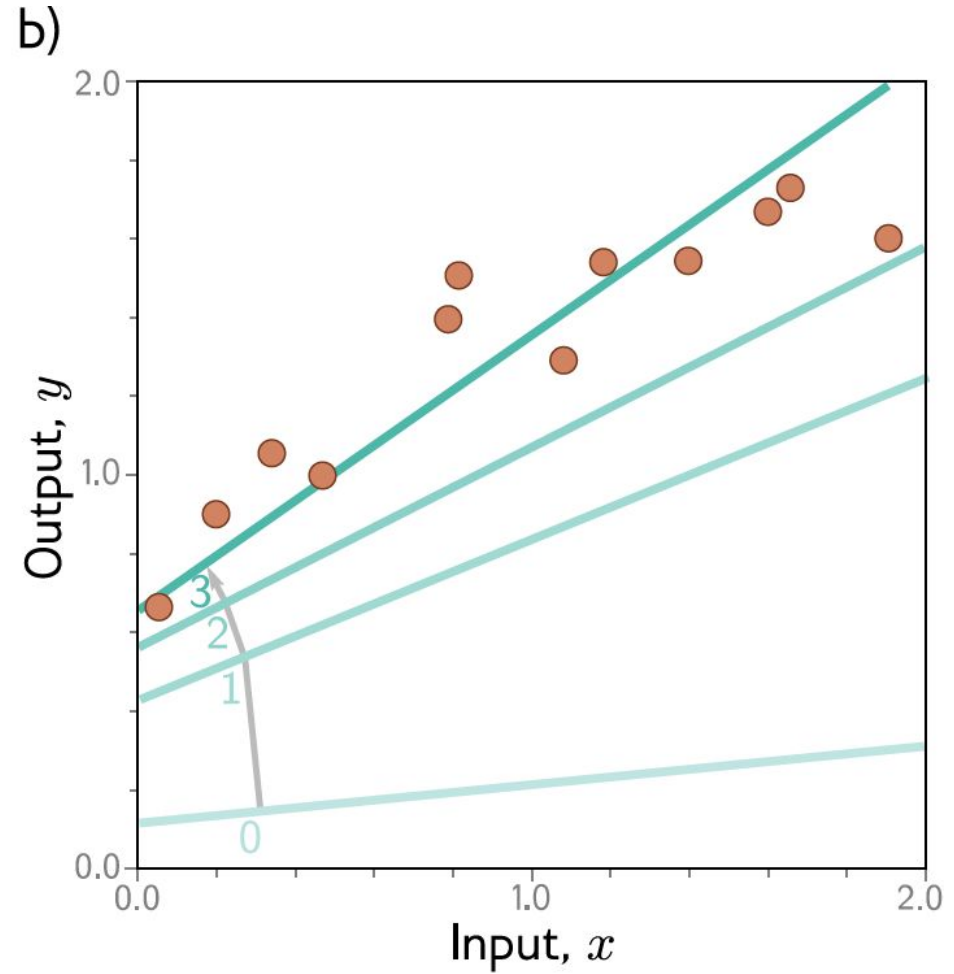

Example: 1D Linear regression training



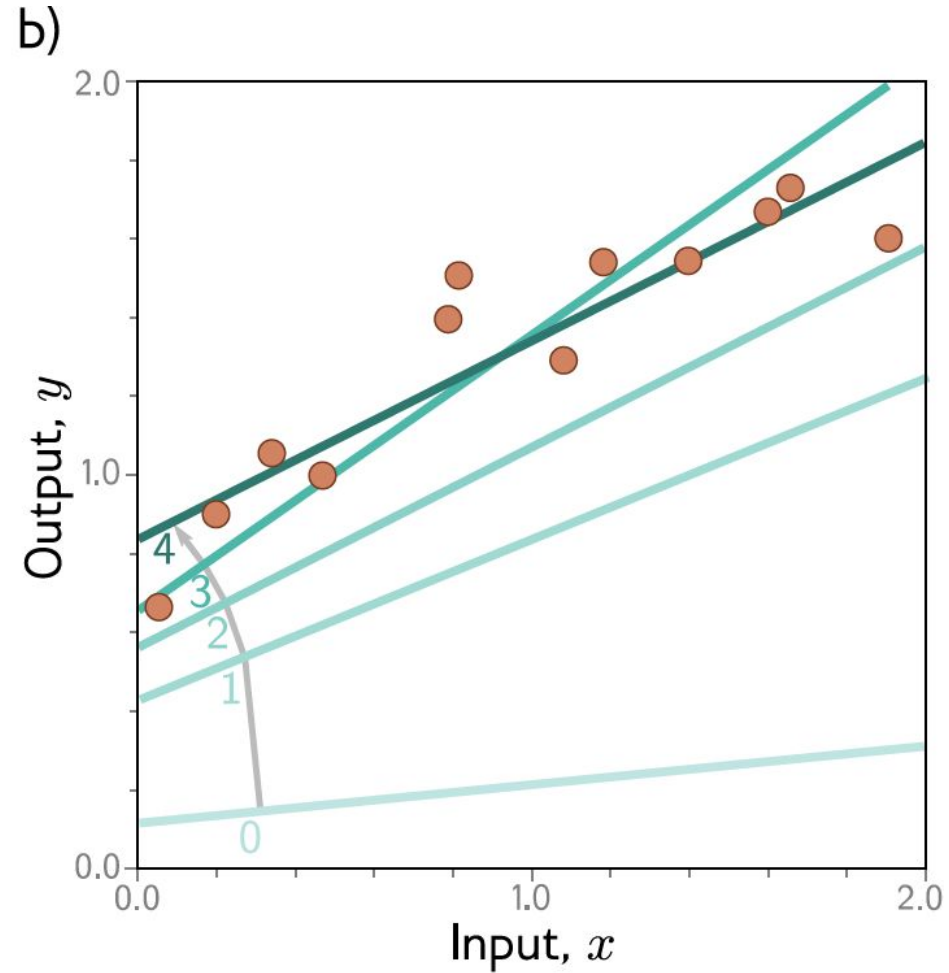

Example: 1D Linear regression training



Example: 1D Linear regression training



Example: 1D Linear regression training



This technique is known as **gradient descent**


gradient descent

Possible objections

- But you can fit the line model in closed form!
 - Yes – but we won't be able to do this for more complex models
- But we could exhaustively try every slope and intercept combo!
 - Yes – but we won't be able to do this when there are millions of (~175 billion in GPT-3) parameters

Example: 1D Linear regression testing


- Test with different set of paired input/output data
 - Measure performance ✓
 - Degree to which this is same as training = generalization }
- Might not generalize well because
 - Model too simple
 - Model too complex
 - fits to statistical peculiarities of data
 - this is known as overfitting ✓



Supervised learning

- Overview
- Notation
 - Model
 - Loss function
 - Training
 - Testing
- 1D Linear regression example
 - Model
 - Loss function
 - Training
 - Testing
- Where are we going?

Where are we going?


$$\theta_0 + \theta_1 x$$

- Shallow neural networks (a more flexible model)
 - Deep neural networks (an even more flexible model)
 - Loss functions (where did least squares come from?)
 - How to train neural networks (gradient descent and variants)
 - How to measure performance of neural networks (generalization)
- 